

# 大數據、機器學習與程式交易 之方法實作期末報告

## VIN ROUGE

第九組

財管三 傅玉婷

財管三 詹雅智

財管三 陳尚格

風管三 許佳勛

風管三 歐予揚

風管三 龔 磊

## 紅酒中成分簡介

- » 酒精：由糖份發酵後所得，為葡萄酒增添了芳醇的味道
- » 殘糖：紅酒中甜度的來源，即是發酵結束後紅酒中剩餘的糖分
- » 硫酸鹽：作為防腐劑，防止紅酒變質
- » 固定酸度：葡萄中的酒石酸及蘋果酸屬於固定酸，用以保持紅酒的化學穩定性和色澤
- » 揮發性酸度：檸檬酸
- » PH值：大多數紅酒的pH值介於2.9至3.9之間，影響紅酒的酸度

## 紅酒中成分簡介

- » 游離二氧化硫：防止氧氣造成紅酒的衰敗及氧化
- » 總二氧化硫
- » 氯化物
- » 密度

## 研究問題

- » 利用紅酒的成分評斷紅酒的品質

# Setup

```
In [2]: import pandas as pd

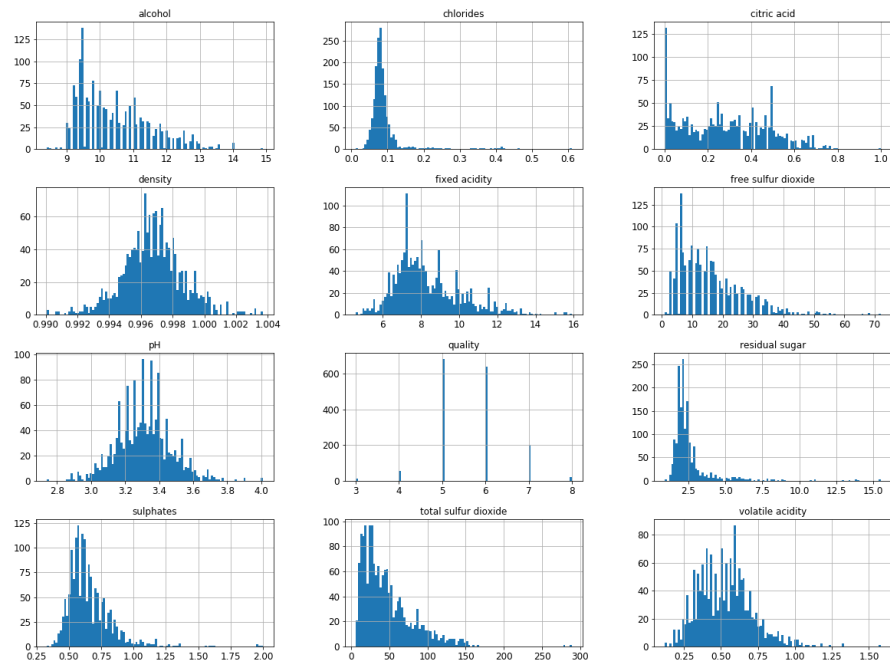
def load_csv_data(file):
    return pd.read_csv(file)

wine = load_csv_data("winequality-red.csv")
wine.head()
```

Out[2]:

	fixed acidity	volatile acidity	citric acid	residual sugar	chlorides	free sulfur dioxide	total sulfur dioxide	density	pH	sulphates	alcohol	quality
0	7.4	0.70	0.00	1.9	0.076	11.0	34.0	0.9978	3.51	0.56	9.4	5
1	7.8	0.88	0.00	2.6	0.098	25.0	67.0	0.9968	3.20	0.68	9.8	5
2	7.8	0.76	0.04	2.3	0.092	15.0	54.0	0.9970	3.26	0.65	9.8	5
3	11.2	0.28	0.56	1.9	0.075	17.0	60.0	0.9980	3.16	0.58	9.8	6
4	7.4	0.70	0.00	1.9	0.076	11.0	34.0	0.9978	3.51	0.56	9.4	5

```
In [5]: %matplotlib inline
import matplotlib.pyplot as plt
wine.hist(bins=100, figsize=(20,15))
plt.savefig("attribute_histogram_plots")
plt.show()
```



```
In [7]: import numpy as np

# For illustration only. Sklearn has train_test_split()
def split_train_test(data, test_ratio):
    shuffled_indices = np.random.permutation(len(data))
    test_set_size = int(len(data) * test_ratio)
    test_indices = shuffled_indices[:test_set_size]
    train_indices = shuffled_indices[test_set_size:]
    return data.iloc[train_indices], data.iloc[test_indices]
```

```
In [8]: train_set, test_set = split_train_test(wine, 0.2)
print(len(train_set), "train +", len(test_set), "test")
```

1280 train + 319 test



```
In [12]: from sklearn.model_selection import train_test_split  
train_set, test_set = train_test_split(wine, test_size=0.2, random_state=42)
```

```
In [13]: test_set.head()
```

```
Out[13]:
```

	fixed acidity	volatile acidity	citric acid	residual sugar	chlorides	free sulfur dioxide	total sulfur dioxide	density	pH	sulphates	alcohol	quality
803	7.7	0.56	0.08	2.50	0.114	14.0	46.0	0.9971	3.24	0.66	9.6	6
124	7.8	0.50	0.17	1.60	0.082	21.0	102.0	0.9960	3.39	0.48	9.5	5
350	10.7	0.67	0.22	2.70	0.107	17.0	34.0	1.0004	3.28	0.98	9.9	6
682	8.5	0.46	0.31	2.25	0.078	32.0	58.0	0.9980	3.33	0.54	9.8	5
1326	6.7	0.46	0.24	1.70	0.077	18.0	34.0	0.9948	3.39	0.60	10.6	6

```
In [14]: wine= train_set.copy()
```

```
In [15]: corr_matrix = wine.corr()  
corr_matrix["quality"].sort_values(ascending=False)
```

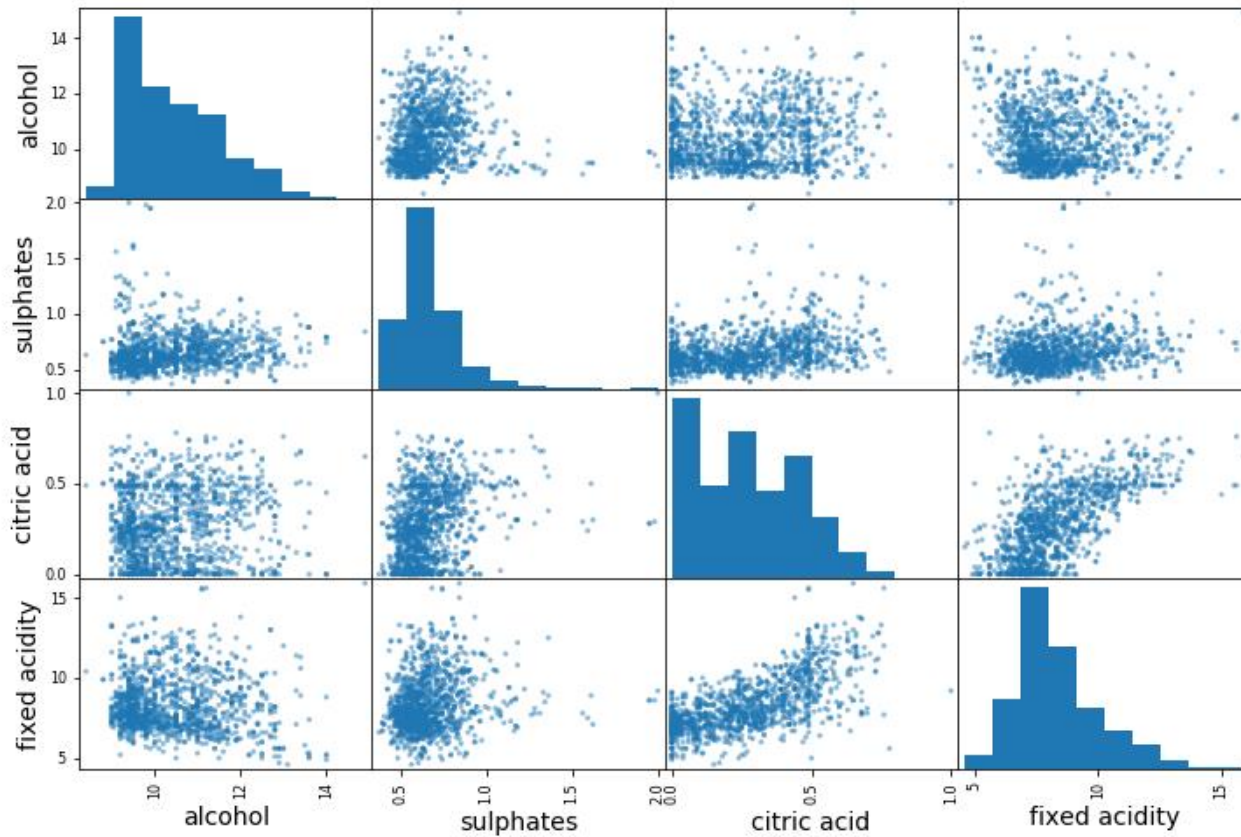
```
Out[15]: quality          1.000000  
alcohol          0.472676  
sulphates        0.242596  
citric acid       0.216115  
fixed acidity     0.122488  
residual sugar    0.005425  
pH               -0.045185  
free sulfur dioxide -0.055860  
chlorides         -0.126541  
density          -0.167091  
total sulfur dioxide -0.200067  
volatile acidity  -0.378372  
Name: quality, dtype: float64
```

```
In [16]: wine["sulfur_ratio"] = wine["free sulfur dioxide"]/wine["total sulfur dioxide"]
wine['acidity_ratio'] = wine["fixed acidity"]/wine["volatile acidity"]
wine["c"] = wine["pH"]/wine["alcohol"]
wine["sulphates_volatile_acidity_ratio"] = wine["sulphates"]/wine["volatile acidity"]
wine["citric_acid_chlorides_ratio"] = wine["citric acid"]/wine["chlorides"]
wine["alcohol_chlorides_ratio"] = wine["alcohol"]/wine["chlorides"]
```

```
In [17]: corr_matrix = wine.corr()
corr_matrix["quality"].sort_values(ascendi
```

```
Out[17]: quality      1.000000
         alcohol      0.472676
         sulphates_volatile_acidity_ratio  0.382378
         acidity_ratio  0.336420
         citric_acid_chlorides_ratio      0.286235
         alcohol_chlorides_ratio          0.274968
         sulphates      0.242596
         citric acid      0.216115
         sulfur_ratio      0.202364
         fixed acidity  0.122488
         residual sugar  0.005425
         pH           -0.045185
         free sulfur dioxide -0.055860
         chlorides      -0.126541
         density        -0.167091
         total sulfur dioxide -0.200067
         volatile acidity -0.378372
         c              -0.486091
         Name: quality, dtype: float64
```

In [



ty"]

```
In [20]: wine = train_set.drop("quality", axis=1) # drop labels for training set
         wine_labels = train_set["quality"].copy()
```

```
In [21]: sample_incomplete_rows = wine[wine.isnull().any(axis=1)].head()
         sample_incomplete_rows
```

Out[21]:

fixed acidity	volatile acidity	citric acid	residual sugar	chlorides	free sulfur dioxide	total sulfur dioxide	density	pH	sulphates	alcohol
---------------	------------------	-------------	----------------	-----------	---------------------	----------------------	---------	----	-----------	---------

```
In [22]: try:
         from sklearn.impute import SimpleImputer # Scikit-Learn 0.20+
         except ImportError:
         from sklearn.preprocessing import Imputer as SimpleImputer

         imputer = SimpleImputer(strategy="median")
```

```
In [23]: imputer.fit(wine)
```

Out[23]: Imputer(axis=0, copy=True, missing\_values='NaN', strategy='median', verbose=0)

```
In [24]: X = imputer.transform(wine)
wine_tr = pd.DataFrame(X, columns=wine.columns,
                        index = list(wine.index.values))
```

```
In [25]: wine_tr.loc[sample_incomplete_rows.index.values]
```

```
Out[25]:
```

	fixed acidity	volatile acidity	citric acid	residual sugar	chlorides	free sulfur dioxide	total sulfur dioxide	density	pH	sulphates	alcohol
--	---------------	------------------	-------------	----------------	-----------	---------------------	----------------------	---------	----	-----------	---------

```
In [26]: wine_tr = pd.DataFrame(X, columns=wine.columns)
wine_tr.head()
```

```
Out[26]:
```

	fixed acidity	volatile acidity	citric acid	residual sugar	chlorides	free sulfur dioxide	total sulfur dioxide	density	pH	sulphates	alcohol
0	8.7	0.690	0.31	3.0	0.086	23.0	81.0	1.00020	3.48	0.74	11.6
1	6.1	0.210	0.40	1.4	0.066	40.5	165.0	0.99120	3.25	0.59	11.9
2	10.9	0.390	0.47	1.8	0.118	6.0	14.0	0.99820	3.30	0.75	9.8
3	8.8	0.685	0.26	1.6	0.088	16.0	23.0	0.99694	3.32	0.47	9.4
4	8.4	1.035	0.15	6.0	0.073	11.0	54.0	0.99900	3.37	0.49	9.9



In [27]: `from sklearn.base import BaseEstimator, TransformerMixin`

*# column index*

`fixed_acidity_ix, volatile_acidity_ix, citric_acid_ix, chlorides_ix, sulphates_ix = 0, 1, 2, 4, 9`

`class CombinedAttributesAdder(BaseEstimator, TransformerMixin):`

`def __init__(self, add_citric_acid_chlorides_ratio = True):`

`self.add_citric_acid_chlorides_ratio = add_citric_acid_chlorides_ratio`

`def fit(self, X, y=None):`

`return self`

`def transform(self, X, y=None):`

`acidity_ratio = X[:, fixed_acidity_ix] / X[:, volatile_acidity_ix]`

`sulphates_volatile_acidity_ratio = X[:, sulphates_ix] / X[:, volatile_acidity_ix]`

`if self.add_citric_acid_chlorides_ratio:`

`citric_acid_chlorides_ratio = X[:, citric_acid_ix] / X[:, chlorides_ix]`

`return np.c_[X, acidity_ratio, sulphates_volatile_acidity_ratio,  
citric_acid_chlorides_ratio]`

`else:`

`return np.c_[X, acidity_ratio, sulphates_volatile_acidity_ratio]`

`attr_adder = CombinedAttributesAdder(add_citric_acid_chlorides_ratio=False)`

`wine_extra_attribs = attr_adder.transform(wine.values)`

```
In [28]: wine_extra_attribs = pd.DataFrame(
    wine_extra_attribs,
    columns=list(wine.columns)+["acidity_ratio", "sulphates_volatile_acidity_ratio"])
wine_extra_attribs.head()
```

Out[28]:

	fixed acidity	volatile acidity	citric acid	residual sugar	chlorides	free sulfur dioxide	total sulfur dioxide	density	pH	sulphates	alcohol	acidity_ratio	sulphates_volatile_acidity_ratio
0	8.7	0.690	0.31	3.0	0.086	23.0	81.0	1.00020	3.48	0.74	11.6	12.608696	1.072464
1	6.1	0.210	0.40	1.4	0.066	40.5	165.0	0.99120	3.25	0.59	11.9	29.047619	2.809524
2	10.9	0.390	0.47	1.8	0.118	6.0	14.0	0.99820	3.30	0.75	9.8	27.948718	1.923077
3	8.8	0.685	0.26	1.6	0.088	16.0	23.0	0.99694	3.32	0.47	9.4	12.846715	0.686131
4	8.4	1.035	0.15	6.0	0.073	11.0	54.0	0.99900	3.37	0.49	9.9	8.115942	0.473430



In [29]:

```
from sklearn.pipeline import Pipeline
from sklearn.preprocessing import StandardScaler

num_pipeline = Pipeline([
    ('imputer', SimpleImputer(strategy="median")),
    ('attrs_adder', CombinedAttributesAdder()),
    ('std_scaler', StandardScaler()),
])

wine_tr = num_pipeline.fit_transform(wine)
```

```
In [30]: wine_tr
```

```
Out[30]: array([[ 0.21833164,  0.88971201,  0.19209222, ..., -0.60592788,  
                 -0.47669384,  0.08693868],  
               [-1.29016623, -1.78878251,  0.65275338, ...,  1.21726382,  
                 1.77377737,  1.00582301],  
               [ 1.49475291, -0.78434707,  1.01104539, ...,  1.095388  ,  
                 0.62532933,  0.2285152  ],  
               ...,  
               [-0.65195559,  0.49909822, -1.08752211, ..., -0.71636823,  
                 -0.7377447  , -0.97018322],  
               [-0.24582155, -1.84458448,  0.39683051, ...,  2.37650605,  
                 3.31611462,  1.16329163],  
               [-1.46422367, -1.34236676, -0.06383064, ...,  0.21381953,  
                 0.54628966,  0.28236718]])
```

```
In [31]: num_attribs = list(wine)  
         wine_prepared = num_pipeline.fit_transform(wine)
```

```
In [32]: wine_prepared.shape
```

```
Out[32]: (1279, 14)
```



# Models

SVR

Random Forest


Linear Regression





# 1. SVR






In [57]: `from sklearn.svm import SVR`

```
svr_reg = SVR(kernel="linear")  
svr_reg.fit(wine_prepared, wine_labels)
```

Out[57]: SVR(C=1.0, cache\_size=200, coef0=0.0, degree=3, epsilon=0.1, gamma='auto',  
kernel='linear', max\_iter=-1, shrinking=True, tol=0.001, verbose=False)

In [58]: `wine_predictions = svr_reg.predict(wine_prepared)`  
`svr_mse = mean_squared_error(wine_labels, wine_predictions)`  
`svr_rmse = np.sqrt(svr_mse)`  
`svr_rmse`

Out[58]: 0.6578738210578761



```
In [59]: final_model = svr_reg

X_test = test_set.drop('quality', axis=1)
y_test = test_set["quality"].copy()

X_test_prepared = num_pipeline.fit_transform(X_test)
final_predictions = final_model.predict(X_test_prepared)

final_mse = mean_squared_error(y_test, final_predictions)
final_rmse = np.sqrt(final_mse)
```

```
In [60]: final_predictions
```

```
Out[60]: array([5.2106567, 4.9463886, 5.62103214, 5.2979142, 5.60017076,
 5.10113022, 4.95252974, 5.0676929, 5.64823137, 5.61821471,
 6.06119103, 5.19622758, 5.38149932, 5.07105862, 5.29468298,
 6.45887704, 5.08451824, 5.49124183, 6.56867462, 5.26608284,
 5.40218388, 5.03450651, 5.93809386, 6.3500509, 5.25271369,
 5.31058552, 6.46757107, 5.2974069, 5.21298527, 6.34559402,
 5.13976537, 5.32273156, 5.60614376, 5.34912516, 5.32820561,
 4.91155736, 6.13345603, 5.6346516, 5.6076249, 6.08225541,
 5.50660038, 5.09294911, 6.13234414, 5.10834518, 5.74931251,
 5.78591251, 6.48398764, 5.5193933, 5.243782, 5.46661775,
 5.23493117, 5.03376404, 5.55756485, 6.38738892, 4.87387565,
 4.85873903, 5.96764009, 5.44490261, 5.73287068, 5.17003608,
 5.42998275, 5.87833925, 5.08210097, 5.18596899, 6.50655778,
 5.30088815, 6.36471818, 5.2025488, 6.45324016, 5.20847886,
 6.41434991, 4.77568387, 5.66931051, 5.70798781, 6.14013769,
 5.18989095, 6.82045528, 5.8225108, 6.15204361, 6.55227946,
```

```
In [61]: final_rmse
```

```
Out[61]: 0.6299629040488907
```

# 交叉驗證 (Cross Validation)

```
In [56]: from sklearn.model_selection import cross_val_score  
svr_scores = cross_val_score(svr_reg, wine_prepared, wine_labels,  
                             scoring="neg_mean_squared_error", cv=10)  
svr_rmse_scores = np.sqrt(-svr_scores)  
display_scores(svr_rmse_scores)
```

Scores: [0.6357058 0.72804982 0.66052214 0.62053514 0.67279226 0.62327226  
0.62173735 0.67546618 0.50332918 0.61555172]

Mean: 0.6356961855178536

Standard deviation: 0.05532556101518299

# 網格搜索 (Grid Search)

```
In [43]: from sklearn.model_selection import GridSearchCV

parameters = {'kernel':('linear','rbf'), 'C':[1, 10, 100], 'gamma':[0.125, 0.25, 0.5, 1, 2, 4]}

svm_reg = SVR(kernel="linear")

grid_search = GridSearchCV(svm_reg, parameters, cv=5,
                           scoring='neg_mean_squared_error', return_train_score=True)
grid_search.fit(wine_prepared, wine_labels)
```

```
Out[43]: GridSearchCV(cv=5, error_score='raise',
                      estimator=SVR(C=1.0, cache_size=200, coef0=0.0, degree=3, epsilon=0.1, gamma='auto',
                      kernel='linear', max_iter=-1, shrinking=True, tol=0.001, verbose=False),
                      fit_params=None, iid=True, n_jobs=1,
                      param_grid={'kernel': ('linear', 'rbf'), 'C': [1, 10, 100], 'gamma': [0.125, 0.25, 0.5, 1, 2, 4]},
                      pre_dispatch='2*n_jobs', refit=True, return_train_score=True,
                      scoring='neg_mean_squared_error', verbose=0)
```

```
In [44]: grid_search.best_estimator_
```

```
Out[44]: SVR(C=1, cache_size=200, coef0=0.0, degree=3, epsilon=0.1, gamma=0.125,
             kernel='rbf', max_iter=-1, shrinking=True, tol=0.001, verbose=False)
```



## 網格搜索 (Grid Search)

```
In [66]: from sklearn.svm import SVR

Svr_reg = grid_search.best_estimator_
Svr_reg.fit(wine_prepared, wine_labels)
wine_predictions = Svr_reg.predict(wine_prepared)
Svr_mse = mean_squared_error(wine_labels, wine_predictions)
Svr_rmse = np.sqrt(svr_mse)
Svr_rmse
```

Out[66]: 0.509766885671461

# 網格搜索 ( Grid Search )

```
In [69]: final_model = grid_search.best_estimator_
```

```
X_test = test_set.drop('quality', axis=1)  
y_test = test_set["quality"].copy()
```

```
X_test_prepared = num_pipeline.fit_transform(X_test)  
final_predictions = final_model.predict(X_test_prepared)
```


```
final_mse = mean_squared_error(y_test, final_predictions)  
final_rmse = np.sqrt(final_mse)
```

```
In [70]: final_predictions
```

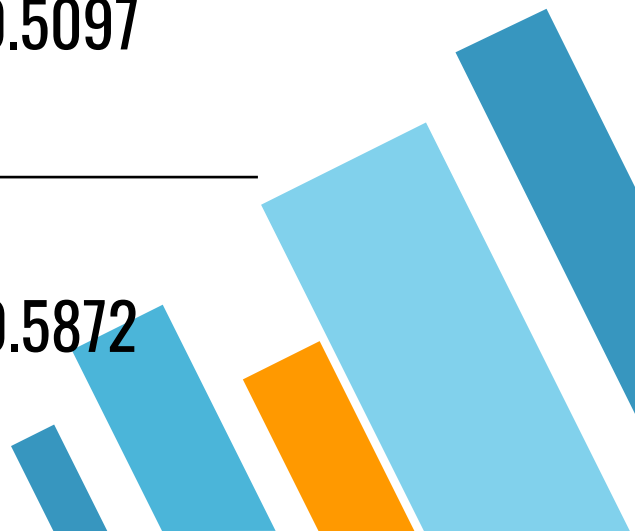
```
Out[70]: array([5.4910721 , 5.0543006 , 5.72279344, 5.31923231, 5.75131126,  
 5.05902878, 4.92111104, 4.91804081, 5.82860687, 5.61504456,  
 6.74146618, 5.1265742 , 5.73685687, 5.15448221, 5.38842549,  
 6.46396182, 5.12526614, 5.54498761, 6.88510682, 4.92623107,  
 4.92940341, 5.22442316, 5.29512663, 6.02383691, 5.47161376,  
 5.93185614, 6.16095167, 5.18741448, 5.08934549, 6.17131731,  
 5.09990855, 5.1460432 , 5.66539649, 5.20316388, 5.74983151,  
 5.03119235, 6.40936384, 5.85403264, 5.55520372, 6.11958158,  
 5.5608377 , 5.17133468, 6.25739337, 5.0746152 , 5.80067184,  
 5.90581548, 6.73530586, 5.69148242, 5.1164224 , 5.58912592,  
 5.16672494, 5.12320866, 5.83158914, 6.9765507 , 4.95014995,  
 5.07098094, 5.95227187, 5.95806407, 5.88498128, 5.13313513,  
 5.67048659, 6.04301937, 5.0967085 , 5.1236814 , 6.45262897,  
 5.48477202, 6.59455097, 5.53495927, 6.63510306, 5.31524714,  
 6.12883453, 5.01628839, 5.71446368, 5.7367934 , 5.99711913,  
 5.00514539, 6.70616729, 5.01535802, 5.57648045, 6.28223092,  
 5.16005775, 6.77117832, 5.30596628, 5.53603404, 6.22900968,  
 6.13762288, 5.02014911, 6.09282268, 6.39110032, 5.02982384,  
 6.19543224, 5.34289422, 5.04611402, 5.46782739, 5.36337818.
```

```
In [71]: final_rmse
```

```
Out[71]: 0.5872263148220179
```



	Grid search 前	Grid search 後
svr_rmse	0.6578	0.5097
final_rmse	0.6356	0.5872





# 2. Random Forest



## 一般情形

In [42]: `from sklearn.ensemble import RandomForestRegressor`

```
forest_reg = RandomForestRegressor(n_estimators=10, random_state=42)
forest_reg.fit(wine_prepared, wine_labels)
```

Out[42]: `RandomForestRegressor(bootstrap=True, criterion='mse', max_depth=None, max_features='auto', max_leaf_nodes=None, min_impurity_decrease=0.0, min_impurity_split=None, min_samples_leaf=1, min_samples_split=2, min_weight_fraction_leaf=0.0, n_estimators=10, n_jobs=1, oob_score=False, random_state=42, verbose=0, warm_start=False)`

In [43]: `wine_predictions = forest_reg.predict(wine_prepared)`  
`forest_mse = mean_squared_error(wine_labels, wine_predictions)`  
`forest_rmse = np.sqrt(forest_mse)`  
`forest_rmse`

Out[43]: 0.24371698168540737

# 交叉驗證 (Cross Validation)

In [47]: `from sklearn.model_selection import cross_val_score`

```
forest_scores = cross_val_score(forest_reg, wine_prepared, wine_labels,  
                                scoring="neg_mean_squared_error", cv=10)  
forest_rmse_scores = np.sqrt(-forest_scores)  
display_scores(forest_rmse_scores)
```

```
Scores: [0.64086124 0.69619502 0.63854082 0.58229073 0.62998016 0.65293807  
         0.553328   0.70055781 0.47145055 0.59167452]  
Mean: 0.6157816923648732  
Standard deviation: 0.06543484441794419
```

# 網格搜索 (Grid Search)

```
In [49]: from sklearn.model_selection import GridSearchCV

param_grid = [
    # try 12 (3x4) combinations of hyperparameters
    {'n_estimators': [3, 10, 30], 'max_features': [2, 4, 6, 8]},
    # then try 6 (2x3) combinations with bootstrap set as False
    {'bootstrap': [False], 'n_estimators': [3, 10], 'max_features': [2, 3, 4]},
]

forest_reg = RandomForestRegressor(random_state=42)
# train across 5 folds, that's a total of (12+6)*5=90 rounds of training
grid_search = GridSearchCV(forest_reg, param_grid, cv=5,
                           scoring='neg_mean_squared_error', return_train_score=True)
grid_search.fit(wine_prepared, wine_labels)
```

```
Out[49]: GridSearchCV(cv=5, error_score='raise',
                      estimator=RandomForestRegressor(bootstrap=True, criterion='mse', max_depth=None,
                                                         max_features='auto', max_leaf_nodes=None,
                                                         min_impurity_decrease=0.0, min_impurity_split=None,
                                                         min_samples_leaf=1, min_samples_split=2,
                                                         min_weight_fraction_leaf=0.0, n_estimators=10, n_jobs=1,
                                                         oob_score=False, random_state=42, verbose=0, warm_start=False),
                      fit_params=None, iid=True, n_jobs=1,
                      param_grid=[{'n_estimators': [3, 10, 30], 'max_features': [2, 4, 6, 8]}, {'bootstrap': [False], 'n_estimators': [3, 10], 'max_features': [2, 3, 4]}],
                      pre_dispatch='2*n_jobs', refit=True, return_train_score=True,
                      scoring='neg_mean_squared_error', verbose=0)
```

## 網格搜索 (Grid Search)

```
In [51]: cvres = grid_search.cv_results_  
for mean_score, params in zip(cvres["mean_test_score"], cvres["params"]):  
    print(np.sqrt(-mean_score), params)  
  
0.7192574284434768 {'max_features': 2, 'n_estimators': 3}  
0.6373260697285106 {'max_features': 2, 'n_estimators': 10}  
0.6226427181605234 {'max_features': 2, 'n_estimators': 30}  
0.6674480695525865 {'max_features': 4, 'n_estimators': 3}  
0.6147839809533946 {'max_features': 4, 'n_estimators': 10}  
0.6094262872716905 {'max_features': 4, 'n_estimators': 30}  
0.7007820445006676 {'max_features': 6, 'n_estimators': 3}  
0.6337460762262236 {'max_features': 6, 'n_estimators': 10}  
0.6147218026250044 {'max_features': 6, 'n_estimators': 30}  
0.7081809690642888 {'max_features': 8, 'n_estimators': 3}  
0.6247376306217943 {'max_features': 8, 'n_estimators': 10}  
0.6072184724658798 {'max_features': 8, 'n_estimators': 30}  
0.6895978036105577 {'bootstrap': False, 'max_features': 2, 'n_estimators': 3}  
0.6293511557531692 {'bootstrap': False, 'max_features': 2, 'n_estimators': 10}  
0.6755982364928798 {'bootstrap': False, 'max_features': 3, 'n_estimators': 3}  
0.6212235398580833 {'bootstrap': False, 'max_features': 3, 'n_estimators': 10}  
0.6871368620247045 {'bootstrap': False, 'max_features': 4, 'n_estimators': 3}  
0.629326308694225 {'bootstrap': False, 'max_features': 4, 'n_estimators': 10}
```



# 網格搜索 ( Grid Search )

```
In [54]: final_model = grid_search.best_estimator_

X_test = test_set.drop('quality', axis=1)
y_test = test_set["quality"].copy()

X_test_prepared = num_pipeline.fit_transform(X_test)
final_predictions = final_model.predict(X_test_prepared)

final_mse = mean_squared_error(y_test, final_predictions)
final_rmse = np.sqrt(final_mse)
```

```
In [56]: final_predictions
```

```
Out[56]: array([5.3      , 5.13333333, 5.63333333, 5.13333333, 5.96666667,
        5.1      , 5.1      , 4.5      , 6.      , 5.96666667,
        6.53333333, 5.23333333, 5.63333333, 5.1      , 5.36666667,
        6.43333333, 5.33333333, 5.7      , 6.6      , 5.03333333,
        4.86666667, 5.5      , 5.33333333, 6.      , 5.56666667,
        5.83333333, 6.43333333, 5.4      , 5.1      , 6.16666667,
        5.23333333, 5.36666667, 5.73333333, 5.46666667, 5.63333333,
        5.13333333, 6.26666667, 6.16666667, 5.26666667, 5.56666667,
        5.2      , 5.23333333, 6.36666667, 5.1      , 5.63333333,
        5.56666667, 6.6      , 5.43333333, 5.23333333, 5.63333333,
        5.03333333, 5.2      , 5.9      , 6.8      , 5.13333333,
        5.2      , 6.      , 5.9      , 5.46666667, 5.26666667,
        5.76666667, 5.96666667, 5.3      , 5.16666667, 6.73333333,
        5.43333333, 6.56666667, 5.43333333, 6.53333333, 5.36666667,
        5.96666667, 5.06666667, 5.8      , 5.5      , 6.06666667,
        4.96666667, 6.63333333, 5.33333333, 5.86666667, 6.5      ,
        5.26666667, 6.83333333, 5.16666667, 5.26666667, 5.66666667,
        6.23333333, 5.06666667, 5.83333333, 6.4      , 5.3      ,
        6.13333333, 5.53333333, 4.96666667, 5.13333333, 5.33333333,
        5.4      , 5.06666667, 5.96666667, 4.66666667, 5.53333333,
        5.06666667, 5.03333333, 5.73333333, 6.36666667, 5.5      ,
        6.26666667, 5.76666667, 5.2      , 5.26666667, 5.23333333])
```

```
In [57]: final_rmse
```

```
Out[57]: 0.5872990199965338
```

# 3. Linear Regression

```
In [31]: from sklearn.linear_model import LinearRegression
```

```
lin_reg = LinearRegression()  
lin_reg.fit(wine_prepared, wine_labels)
```

```
Out[31]: LinearRegression(copy_X=True, fit_intercept=True, n_jobs=1, normalize=False)
```

```
In [32]: # let's try the full preprocessing pipeline on a few training instances
```

```
some_data = wine.iloc[:5]  
some_labels = wine_labels.iloc[:5]  
some_data_prepared = num_pipeline.transform(some_data)  
  
print("Predictions:", lin_reg.predict(some_data_prepared))  
print("Labels:", list(some_labels))
```

```
Predictions: [5.70248099 6.04222528 5.67103903 5.10814467 4.80454583]
```

```
Labels: [6, 6, 6, 5, 5]
```

```
In [33]: from sklearn.metrics import mean_squared_error

         wine_predictions = lin_reg.predict(wine_prepared)
         lin_mse = mean_squared_error(wine_labels, wine_predictions)
         lin_rmse = np.sqrt(lin_mse)
         lin_rmse
```

```
Out[33]: 0.6510118225404404
```

```
In [34]: from sklearn.metrics import mean_absolute_error

         lin_mae = mean_absolute_error(wine_labels, wine_predictions)
         lin_mae
```

```
Out[34]: 0.49947401745314934
```

```
In [35]: final_model = lin_reg

X_test = test_set.drop('quality', axis=1)
y_test = test_set["quality"].copy()

X_test_prepared = num_pipeline.fit_transform(X_test)
final_predictions = final_model.predict(X_test_prepared)

final_mse = mean_squared_error(y_test, final_predictions)
final_rmse = np.sqrt(final_mse)
```

```
In [36]: final_predictions
```

```
Out[36]: array([5.32847131, 5.00996722, 5.69223701, 5.43005255, 5.68472085,
5.23970017, 4.99290139, 5.07585415, 5.77593875, 5.66067669,
6.04990652, 5.22001009, 5.53605794, 5.23948605, 5.41812559,
6.40602976, 5.10661285, 5.5567362 , 6.54525118, 5.32687526,
5.32211035, 5.18472542, 5.88328741, 6.28922889, 5.33364881,
5.43610721, 6.35465659, 5.30762198, 5.1771224 , 6.16721335,
5.21900716, 5.44698054, 5.75605374, 5.39175963, 5.45144927,
4.98866114, 6.16325005, 5.68453196, 5.62780589, 6.12707259,
5.54316951, 5.21287501, 6.2057751 , 5.14305715, 5.85195664,
5.83666934, 6.44930838, 5.63674665, 5.11688107, 5.52616358,
5.20132329, 5.04784888, 5.53915679, 6.31006995, 4.93988594,
4.96381788, 5.99675206, 5.38908488, 5.79021138, 5.24083139,
5.57771606, 5.93032765, 5.24207949, 5.28282488, 6.43851765,
5.38827815, 6.29335715, 5.27667348, 6.45026079, 5.29709479,
6.34964133, 4.7203089 , 5.80460326, 5.76818275, 6.12141758,
5.27593295, 6.69248419, 5.85727711, 6.05895817, 6.411229 ,
5.27800801, 6.42487982, 5.43767391, 5.6540473 , 5.68869137,
6.38372194, 5.27789988, 5.82504077, 6.30636083, 5.19429432,
6.04799061, 5.64150782, 5.74890402, 5.89092357, 5.16986564,
5.73563468, 5.13670626, 5.68928997, 4.95109922, 5.48801781,
5.04572503, 5.11917556, 5.78248621, 5.6930024 , 5.42420134,
6.11899545, 5.72823355, 5.40563171, 6.03988128, 5.2256793 ,
6.62534397, 5.22454458, 6.1786867 , 4.67707492, 5.8100759 ,
5.97236542, 6.07865402, 5.47341996, 4.99015957, 5.78068566,
6.15156422, 5.29367349, 5.77121159, 5.36002283, 5.38789113,
5.82251214, 5.81682226, 5.66252222, 5.51521522, 5.75622222])
```

```
In [37]: final_rmse
```

```
Out[37]: 0.6231238739209931
```

# 交叉驗證 (Cross Validation)

```
In [38]: from sklearn.model_selection import cross_val_score

def display_scores(scores):
    print("Scores:", scores)
    print("Mean:", scores.mean())
    print("Standard deviation:", scores.std())

lin_scores = cross_val_score(lin_reg, wine_prepared, wine_labels,
                             scoring="neg_mean_squared_error", cv=10)
lin_rmse_scores = np.sqrt(-lin_scores)
display_scores(lin_rmse_scores)

Scores: [0.62960105 0.71841434 0.69363621 0.70150314 0.68016265 0.67488408
 0.66711484 0.68068173 0.4986116  0.64727016]
Mean: 0.659187980463089
Standard deviation: 0.05873527951600096
```

## 網格搜索 (Grid Search)

```
In [39]: from sklearn.model_selection import GridSearchCV
param_grid = { 'fit_intercept' : [True,False] , 'normalize' : [ True,False]}
grid_search = GridSearchCV(lin_reg, param_grid, cv=5,
                           scoring='neg_mean_squared_error', return_train_score=True)
grid_search.fit(wine_prepared, wine_labels)
```

```
Out[39]: GridSearchCV(cv=5, error_score='raise',
                      estimator=LinearRegression(copy_X=True, fit_intercept=True, n_jobs=1, normalize=False),
                      fit_params=None, iid=True, n_jobs=1,
                      param_grid={'fit_intercept': [True, False], 'normalize': [True, False]},
                      pre_dispatch='2*n_jobs', refit=True, return_train_score=True,
                      scoring='neg_mean_squared_error', verbose=0)
```

```
In [40]: grid_search.best_estimator_
```

```
Out[40]: LinearRegression(copy_X=True, fit_intercept=True, n_jobs=1, normalize=True)
```

# 網格搜索 (Grid Search)

```
In [41]: cvres = grid_search.cv_results_
for mean_score, params in zip(cvres["mean_test_score"], cvres["params"]):
    print(np.sqrt(-mean_score), params)
```

```
0.6643025610657023 {'fit_intercept': True, 'normalize': True}
0.6643025610657023 {'fit_intercept': True, 'normalize': False}
5.7301986610454705 {'fit_intercept': False, 'normalize': True}
5.7301986610454705 {'fit_intercept': False, 'normalize': False}
```

```
In [42]: pd.DataFrame(grid_search.cv_results_)
```

Out[42]:

	mean_fit_time	std_fit_time	mean_score_time	std_score_time	param_fit_intercept	param_normalize	params	split0_test_score	split1_test_score	split2_
0	0.005258	0.004755	0.000835	0.000526	True	True	{'fit_intercept': True, 'normalize': True}	-0.463220	-0.485396	
1	0.002153	0.001246	0.000324	0.000121	True	False	{'fit_intercept': True, 'normalize': False}	-0.463220	-0.485396	
2	0.001336	0.000395	0.000298	0.000086	False	True	{'fit_intercept': False, 'normalize': True}	-32.912122	-32.330829	
3	0.001067	0.000082	0.000256	0.000011	False	False	{'fit_intercept': False, 'normalize': False}	-32.912122	-32.330829	

4 rows x 22 columns



## 網格搜索 (Grid Search)

```
In [43]: from sklearn.linear_model import LinearRegression

lin_reg = LinearRegression(copy_X=True, fit_intercept=True, n_jobs=1, normalize=True)
lin_reg.fit(wine_prepared, wine_labels)
wine_predictions = lin_reg.predict(wine_prepared)
lin_mse = mean_squared_error(wine_labels, wine_predictions)
lin_rmse = np.sqrt(lin_mse)
lin_rmse
```

```
Out[43]: 0.6510118225404404
```

# 網格搜索 (Grid Search)


```
In [44]: final_model = grid_search.best_estimator_  
  
X_test = test_set.drop('quality', axis=1)  
y_test = test_set["quality"].copy()  
  
X_test_prepared = num_pipeline.fit_transform(X_test)  
final_predictions = final_model.predict(X_test_prepared)  
  
final_mse = mean_squared_error(y_test, final_predictions)  
final_rmse = np.sqrt(final_mse)
```

```
In [45]: final_predictions
```

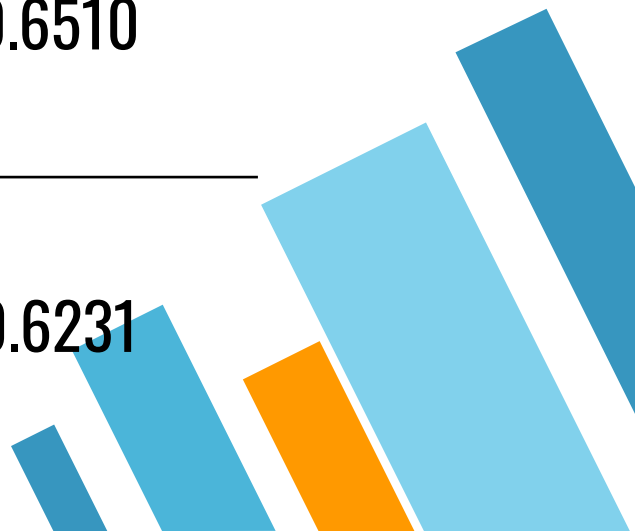
```
Out[45]: array([5.32847131, 5.00996722, 5.69223701, 5.43005255, 5.68472085,  
                5.23970017, 4.99290139, 5.07585415, 5.77593875, 5.66067669,  
                6.04990652, 5.22001009, 5.53605794, 5.23948605, 5.41812559,  
                6.40602976, 5.10661285, 5.5567362 , 6.54525118, 5.32687526,  
                5.32211035, 5.18472542, 5.88328741, 6.28922889, 5.33364881,  
                5.43610721, 6.35465659, 5.30762198, 5.1771224 , 6.16721335,  
                5.21900716, 5.44698054, 5.75605374, 5.39175963, 5.45144927,  
                4.98866114, 6.16325005, 5.68453196, 5.62780589, 6.12707259,  
                5.54316951, 5.21287501, 6.2057751 , 5.14305715, 5.85195664,  
                5.83666934, 6.44930838, 5.63674665, 5.11688107, 5.52616358,  
                5.20132329, 5.04784888, 5.53915679, 6.31006995, 4.93988594,  
                4.96381788, 5.99675206, 5.38908488, 5.79021138, 5.24083139,  
                5.57771606, 5.93032765, 5.24207949, 5.28282488, 6.43851765,  
                5.38827815, 6.29335715, 5.27667348, 6.45026079, 5.29709479,  
                6.34964133, 4.7203089 , 5.80460326, 5.76818275, 6.12141758,  
                5.27593295, 6.69248419, 5.85727711, 6.05895817, 6.411229 ,  
                5.27800801, 6.42487982, 5.43767391, 5.6540473 , 5.68869137,  
                6.38372194, 5.27789988, 5.82504077, 6.30636083, 5.19429432,  
                6.04799061, 5.64150782, 5.74890402, 5.89092357, 5.16986564,  
                5.73563468, 5.13670626, 5.68928997, 4.95109922, 5.48801781,  
                5.04572503, 5.11917556, 5.78248621, 5.6930024 , 5.42420134,  
                6.11899545, 5.72823355, 5.40563171, 6.03988128, 5.2256793 ,  
                6.62534397, 5.22454458, 6.1786867 , 4.67707492, 5.8100759 ,  
                5.97236542, 6.07865402, 5.47341996, 4.99015957, 5.78068566,  
                6.15156422, 5.29367349, 5.77121159, 5.36002283, 5.38789113,
```

```
In [46]: final_rmse
```

```
Out[46]: 0.6231238739209932
```



	Grid search 前	Grid search 後
lin_rmse	0.6510	0.6510
final_rmse	0.6231	0.6231





SVR	0.5872
Random Forest	0.5873
Linear Regression	0.6231



The background features a serene landscape with layered mountain ranges under a soft, orange-hued sky, suggesting a sunset or sunrise. In the top-left corner, there are several overlapping, semi-transparent light-orange rectangular blocks arranged in a fan-like pattern. In the bottom-right corner, there are several overlapping, semi-transparent light-blue rectangular blocks arranged in a similar fan-like pattern.

感謝聆聽